

CS60050
MACHINE LEARNING
Assignment 2

Animesh Jha - 19CS10070

Nisarg Upadhyaya - 19CS30031

1. Code

There are 3 files -

1. main.py

This file contains code to run the experiments. It has the following methods:

- `read_data()`: this function reads “train.csv” and returns a pandas dataframe frame
- `get_M_matrix()`: this function first counts the number of unique words which are more than two letters long, not stopwords and match the given regular expression. After this it constructs the $r \times c$ binary matrix M , with $M[i][j]=1$ iff j -th word is present in the text of the i -th example
- `train_test_split()`: this generates stratified 70:30 splits on M and author names.
- `run_experiment()`: this runs the experiment for given test train split and laplace correction factor(α)
- `Label_encoder()`: we use this library function to match author names to the whole numbers 0,1,2 for easier predictions

2. model.py

This file contains the class NaiveBayes, it implements the algorithm, it has the following methods

- `__init__()`: constructor of the class, stores α and number of classes, and initializes label counts, `label_word` and `total_label_word` counts to all zeros
- `fit()`: this method fits the naive bayes model to the given train dataset. It computes `label_total_text_counts`, `label_total_word_counts` and `label_word_counts` (Counts how many words per label, the frequency of the word for a label, and number of words for a label)
- `log_p_doc()`: this method gives the log of the probability $P(\text{word}+\alpha | \text{label}+\text{vocab}*\alpha)$ for given word and label. This is used for predictions
- `prior()`: computes the prior for a given label
- `predict()`: predicts the labels for a given test dataset, returns the predictions and the probability of the chosen prediction.

3. train.csv

The given dataset. It has the following attributes

- 1) Text
- 2) Author – Label

Procedure followed

We use the standard Naive Bayes algorithm. One thing to note is that, we compute the log probabilities for numerical stability, it is better to work with the sum of logs than to work with the multiplication of a bunch of small numbers. However when there is no laplace correction, for a few examples the probabilities become 0, to handle this we have defined $\log(0)$ to be a very large negative number (-1000000).

Class mapping:

0 -> EAP

1 -> HPL

2 -> MWS

Total samples: 19579

Vocab size: 24787

Shape X train (13705, 24787)

Shape y train (13705,)

Shape X test (5874, 24787)

Shape y test (5874,)

How to run the code?

After installing all dependencies just run ``python main.py`` and enjoy! README has more details.

2. Experiments

Class mapping:

```
0 -> EAP
1 -> HPL
2 -> MWS
```

Total samples: 19579

Vocab size: 24787

Shape X train (13705, 24787)

Shape y train (13705,)

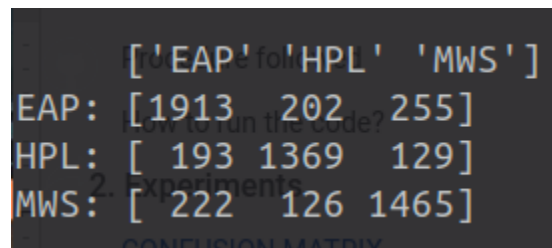
Shape X test (5874, 24787)

Shape y test (5874,)

a. Without Laplace Correction

CONFUSION MATRIX

True values along the rows, predictions along the columns



```
[ 'EAP' 'HPL' 'MWS' ]
EAP: [1913 202 255]
HPL: [ 193 1369 129]
MWS: [ 222 126 1465]
```

Accuracy on test split = 0.8081375553285666

95% confidence interval of accuracy on test split = 0.010069936500308747

Hence, accuracy is between 0.7980676188282578 and 0.8182074918288753

Results for class EAP:

```
True positives: 1913
True negatives: 3089
False positives: 415
False negatives: 457
Sensitivity/Recall: 0.8071729957805908
Specificity: 0.8815639269406392
Precision: 0.8217353951890034
F-score: 0.8143891017454237
```

Results for class HPL:

True positives: 1369
True negatives: 3855
False positives: 328
False negatives: 322
Sensitivity/Recall: 0.8095801301005322
Specificity: 0.9215873774802773
Precision: 0.8067177371832646
F-score: 0.80814639905549

Results for class MWS:

True positives: 1465
True negatives: 3677
False positives: 384
False negatives: 348
Sensitivity/Recall: 0.8080529509100938
Specificity: 0.9054420093573011
Precision: 0.7923201730665225
F-score: 0.8001092299290005

b. With Laplace Correction ($\alpha = 1$)

CONFUSION MATRIX

True values along the rows, predictions along the columns

```
Code
      ['EAP' 'HPL' 'MWS']
EAP: [1944  151  275]
HPL: [ 175 1372  144]
MWS: [ 158   94 1561]
```

Accuracy on test split = 0.8302689819543753

95% confidence interval of accuracy on test split = 0.009600174231490202

Hence, accuracy is between 0.8206688077228851 and 0.8398691561858654

Results for class EAP:

True positives: 1944
True negatives: 3171
False positives: 333
False negatives: 426
Sensitivity/Recall: 0.8202531645569621
Specificity: 0.9049657534246576
Precision: 0.8537549407114624
F-score: 0.8366688185926404

Results for class HPL:

True positives: 1372
True negatives: 3938
False positives: 245
False negatives: 319
Sensitivity/Recall: 0.8113542282672974
Specificity: 0.9414295959837438
Precision: 0.8484848484848485
F-score: 0.8295042321644499

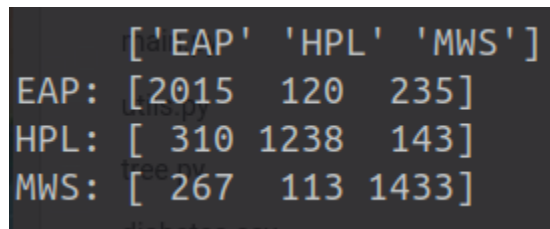
Results for class MWS:

True positives: 1561
True negatives: 3642
False positives: 419
False negatives: 252
Sensitivity/Recall: 0.861003861003861
Specificity: 0.8968234425018469
Precision: 0.7883838383838384
F-score: 0.8230951753229634

c. With Laplace Correction ($\alpha = 10$)

CONFUSION MATRIX

True values along the rows, predictions along the columns



	'EAP'	'HPL'	'MWS'
EAP:	2015	120	235
HPL:	310	1238	143
MWS:	267	113	1433

Accuracy on test split = 0.797752808988764

95% confidence interval of accuracy on test split = 0.010272224929825926

Hence, accuracy is between 0.7874805840589381 and 0.8080250339185899

Results for class EAP:

True positives: 2015
True negatives: 2927
False positives: 577
False negatives: 355
Sensitivity/Recall: 0.8502109704641351
Specificity: 0.8353310502283106
Precision: 0.777391975308642
F-score: 0.8121725110842403

Results for class HPL:

True positives: 1238
True negatives: 3950
False positives: 233
False negatives: 453
Sensitivity/Recall: 0.7321111768184506
Specificity: 0.9442983504661726
Precision: 0.8416043507817811
F-score: 0.7830487033523086

Results for class MWS:

True positives: 1433
True negatives: 3683
False positives: 378
False negatives: 380
Sensitivity/Recall: 0.7904026475455047
Specificity: 0.9069194779610933
Precision: 0.7912755383765875
F-score: 0.7908388520971303

3. Analysis

1. Laplace Correction

As alpha increases, the likelihood probability moves towards uniform distribution.

Since we are not getting much information from that, it is not preferable. Therefore, it is not preferred to use very large values of alpha. Generally, taking $\alpha = 1$ suffices as it smoothens the likelihood while maintaining the original distribution. Hence, the slight reduction in accuracy taking $\alpha = 10$.


Also, on smoothening we can observe a slight increase in the accuracy on the test set. This is because smoothening maintains the likelihood but numerically the values are better distributed (not 0) and hence we get better results.

2. Correlation and Dependence of Features

(Code for this part of the analysis can be found in the file `test_and_corr.py`)

We wanted to test the naive bayes assumption, because the results of the model are quite good. First observe that if two variables are correlated they can not be independent of each other^[1]. Note that zero correlation does not imply independence, instead we focus on the fact that two correlated variables will never be independent.

Therefore to check independence we check correlation of the features. To this end we randomly sample 5% of the dataset, off the 979 samples we get we then count the 20 most commonly appearing words. On this reduced dataset we compute the correlation matrix, we observe that for most part the correlation between two different features is really small (in the figure purple color denotes low values) and close to zero. We also compute the correlation of each word with the label and observe that for most part the correlation is really small. Therefore the data while not completely independent is “nearly” independent which is why the naive bayes model is able to do well even though it makes wrong assumptions about the dataset.



Correlation matrix for the 20 most common words in the randomly sampled 5% dataset

0	-0.063543
9	-0.053651
17	-0.052063
12	-0.016584
6	-0.015318
3	-0.014600
2	-0.010554
10	-0.004828
16	0.010618
5	0.017574
7	0.029308
1	0.033603
4	0.033662
18	0.034168
11	0.041597
19	0.054157
14	0.055277
8	0.079333
13	0.082493
15	0.127957
y	1.000000

Correlation of the 20 most common words with the label

Reference:

[1] Proof specified in

<https://stats.stackexchange.com/questions/113417/does-non-zero-correlation-imply-dependence>